



Convolutional Neural Network (CNN) Mini Project: MNIST Digit Classification

Project Overview

This project implements a Convolutional Neural Network (CNN) using TensorFlow/Keras to accurately classify handwritten digits (0-9) from the popular MNIST dataset. This serves as a foundational "Hello World" deep learning project, demonstrating key concepts of image recognition, data preprocessing, model architecture design, and performance evaluation.

Detail	Value
Problem Type	Multi-class Classification
Dataset	MNIST (Modified National Standards and Technology)
Model	Sequential CNN
Framework	TensorFlow 2.x / Keras
Typical Accuracy	98.5% – 99.4%



How to Run the Project (Google Colab)

This project is designed to be executed step-by-step in a [Google Colaboratory](#) environment, which provides all necessary dependencies (TensorFlow, NumPy, Matplotlib) and free GPU acceleration.

1. **Create a New Colab Notebook:** Go to [Google Colab](#) and start a new notebook.
2. **Paste Code:** Copy the content of the primary notebook file (`MNIST_CNN_Project.md`) and paste it into the first code cell.
3. **Set Runtime:** (Optional but recommended) Go to Runtime > Change runtime type and select GPU as the hardware accelerator.
4. **Execute Cells:** Run each cell sequentially (using `Shift + Enter` or the Run button) from Section 1 to Section 7.

Key Steps and Components

1. Data Loading and Preprocessing

The MNIST dataset is loaded directly from Keras. Critical preprocessing steps include:


- **Normalization:** Scaling pixel values from the initial range of 0 to 255 down to 0.0 to 1.0. This helps the neural network converge faster.
- **Reshaping:** Adding a color channel dimension to transform the images from `(28, 28)` to `(28, 28, 1)`, which is the expected input shape for a `Conv2D` layer.
- **One-Hot Encoding:** Converting integer labels (e.g., `5`) into categorical vectors (e.g., `[0, 0, 0, 0, 0, 1, 0, 0, 0, 0]`) for use with the `categorical_crossentropy` loss function.

2. CNN Model Architecture

The model uses a simple Sequential API structure:

1. `Conv2D(32, (3, 3), activation='relu')` : Learns 32 distinct features from the images.

2. `MaxPooling2D((2, 2))` : Halves the spatial dimensions, reducing computational load and increasing robustness to slight shifts.
3. `Conv2D(64, (3, 3), activation='relu')` : Further feature extraction.
4. `MaxPooling2D((2, 2))` : Second spatial reduction.
5. `Flatten()` : Converts the 2D feature maps into a 1D vector.
6. `Dense(128, activation='relu')` : A hidden layer for high-level classification logic.
7. `Dense(10, activation='softmax')` : The output layer, yielding a probability distribution over the 10 classes (digits 0-9).



3. Training and Evaluation

The model is trained for 8 epochs using the Adam optimizer and Categorical Crossentropy loss. After training, the performance is evaluated on the dedicated test set.



Results and Visualization

The project generates visualizations to assess the model's learning process:

- **Final Test Accuracy:** Reports the final classification accuracy on unseen data.
- **Training History Plots:** Two plots are generated using Matplotlib:
 - **Accuracy vs. Epoch:** Shows how training and validation accuracy changes over each training cycle.
 - **Loss vs. Epoch:** Shows how training and validation loss decreases over each training cycle.
- **Sample Predictions:** Displays a few test images along with the model's prediction, highlighting correct (green) and incorrect (red) classifications.

Expected Performance (Example)

Metric	Value
Test Loss	$\approx 0.03 - 0.06$
Test Accuracy	$\approx 99.00\%$